

# Varieties and Applications of Autoencoders

Isabel Restrepo

Paul Stey

March 17, 2017

# Table of Contents

- 1 Autoencoders
  - Autoencoder Basics
- 2 Stacked Autoencoder
  - Extending Autoencoders
- 3 Sparse Autoencoder
  - k-Sparse Autoencoder
  - Sparsity Constrained Autoencoder
- 4 Denoising Autoencoder
  - Denoising Autoencoder
  - Extending Denoising Autoencoder
- 5 Conclusion
  - Advantages of Autoencoders

# Autoencoder Basics


- 1 Feed-forward neural network
- 2 Used for unsupervised learning
- 3 Trained to reproduce input layer in the output layer
- 4 Training uses gradient descent

# Autoencoder Basics cont.

- ① Similar to PCA\*, but more flexible
- ② Autoencoders can accommodate non-linear transformations
  - 1 Consequence of flexible activation functions
  - 2 Turns out to be hugely useful

\* Actually, with squared error loss and no sigmoid transformation, it is equivalent to PCA (Baldi & Hornik, 1989)

# Simple Autoencoder



# Simple Autoencoder cont.

From input layer to hidden layer is “encoder”:

$$f(x) = \sigma(Wx + b)$$

Hidden layer to the output is the “decoder”:

$$g(x) = \sigma(W'x + b')$$

where  $\sigma$  is the activation function (often sigmoid) and  $W'$  is often  $W^\top$  (known as having “tied” weights).

# Simple Autoencoder cont.

- 1 Input layer  $\rightarrow$  hidden layer (fewer units)  $\rightarrow$  output layer
- 2 Hidden layers values can be thought of as lossy compression of input layer
- 3 Hidden layer with fewer units than input is often called “bottleneck” or under-complete layer

# Question

In a sense, autoencoders want to learn an approximation to the identity function.

Why would this be useful?



# Answer


This turns out to have several practical uses.

- 1 Autoencoders learn a kind of latent representation of the input
- 2 Can be used for initializing weights and biases prior to fitting neural net
- 3 Very well suited to dimensionality reduction in NLP tasks

# Stacked Autoencoder

- 1 Sometimes called deep autoencoder
- 2 Feed-forward neural network
- 3 Has additional layers
- 4 Still “unsupervised” in the sense of no labels or real-valued outcome variable
- 5 Output layer is still  $\hat{x}$
- 6 Trained in stages

# Stacked Autoencoder



# Stacked Autoencoder cont.

- ① Beneficial for data with hierarchical organization
- ② Can be difficult to train using back propagation

# Sparse Autoencoder

- 1 Different approach to “compressing” representation
- 2 Hidden layer has more units than input (i.e., over-complete)
- 3 Compression is therefore achieved through sparsity


# Sparse Autoencoder cont.

- ① Several methods of introducing sparsity, we'll discuss 2
  - ①  $k$ -sparse autoencoders (Makhzani et al., 2013)
  - ② Constrained activation

# $k$ -sparse Autoencoder

- 1 Select  $k$  highest-activation units
- 2 Set others to 0
- 3 Compression is achieved through sparsity of model, not reduced dimensionality

# $k$ -sparse Autoencoder





# Sparsity Constraint

- 1 With sigmoidal activation function,  $\sigma_j(x) \in [0, 1]$
- 2 We want to constrain them to be mostly 0
- 3 Take mean activation

$$\hat{\rho}_j = \frac{1}{n} \sum_{i=1}^n [\sigma_j(x_i)]$$

such that

$$\hat{\rho}_j = \rho$$

where  $\rho$  is a hyperparameter set near 0, commonly 0.05.

# Sparsity Constraint cont.

- 1 Add penalty term to loss function to achieve this sparsity constraint

$$\sum_{j=1}^m \rho \log \frac{\rho}{\hat{\rho}_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j}$$

- 2 This is based on Kullback-Leibler Divergence, and often written as

$$\sum_{j=1}^m \text{KL}(\rho || \hat{\rho}_j)$$

# Sparsity Constraint cont.

- 1 Penalty term has the property that  $\text{KL}(\rho || \hat{\rho}_j) = 0$  when  $\hat{\rho}_j = \rho$
- 2 Otherwise, monotone increasing as  $\hat{\rho}_j$  diverges from  $\rho$
- 3 So our loss function is now

$$L_{\text{sparse}}(x, \hat{x}) = L(x, \hat{x}) + \lambda \sum_{j=1}^m \text{KL}(\rho || \hat{\rho}_j)$$

where  $L(x, \hat{x})$  is our unconstrained loss function and  $\lambda$  controls the weight of sparsity penalty.

# Denoising Autoencoder


- ① Stochastic version of autoencoder
- ② Random noise injected into the input layer
  - 1 *Masking noise*: some fraction of elements\* in  $x$  set to 0
  - 2 *Gaussian additive noise*:  $\tilde{x}|x \sim \mathcal{N}(x, \sigma^2 I)$
  - 3 *Salt-and-pepper noise*: fraction of elements\* in  $x$  set to either min or max possible value (often 0 or 1) according to coin toss

\* chosen at random for each example

# Denoising Autoencoder cont.

- ① Output layer is the original (uncorrupted) data
- ② Goal is slightly different from simple autoencoder
  - 1 Not attempting to learn identity function
  - 2 We want robustness of representation
  - 3 Learned representation is insensitive to perturbations in the input

# Denoising Autoencoder



# Denoising Autoencoder cont.

- ① Training proceeds fairly similarly to standard autoencoder
- ② Crucial difference is the loss function is calculated using *uncorrupted* input layer

# Extending Denoising Autoencoder

- 1 Some noise types only corrupt subset of the input's components
  - 1 Masking noise
  - 2 Salt-and-pepper noise
- 2 For these, we can extend denoising by adding emphasis on corrupted components.
- 3 Use hyperparameters  $\alpha$  and  $\beta$  to control emphasis; for squared error loss, this yields

$$L(x, \hat{x}) = \alpha \left( \sum_{i \in \mathcal{I}(\tilde{x})} (x_i - \hat{x}_i)^2 \right) + \beta \left( \sum_{i \notin \mathcal{I}(\tilde{x})} (x_i - \hat{x}_i)^2 \right)$$

where  $\mathcal{I}(\tilde{x})$  denotes indexes components of  $x$  that were corrupted.



# Some Advantages of Autoencoders

- ① Compressed representation of features used for initializing weights
  - 1 Faster learning
  - 2 Less likely to get stuck in local minima
  - 3 Better approximation → increased accuracy
- ② Well suited to dimensionality reduction in NLP problems
- ③ Allows us to make use of unlabeled data

# Disadvantages of Autoencoders

- 1 Computational burden
- 2 Limitations of unsupervised learning

# Not Discussed

- 1 Winner-Take-All autoencoders (for sparsity)
- 2 Convolutional autoencoders (several varieties)
- 3 Variational autoencoders

# References

- 1 Vincent P., Larochelle, H., Bengio Y., Manzagol P.A. (2008) *Extracting and Composing Robust Features with Denoising Autoencoders.*
- 2 Vincent P., Larochelle, H., LaJoie I., Bengio Y., Manzagol P.A. (2010) *Stacked Denoising Autoencoders: Learning Useful Representations in a Deep Network with a Local Denoising Criterion*
- 3 Makhzani A., Frey B. (2015) *Winner-Take-All Autoencoders*
- 4 Ng A. *Sparse Autoencoders* CS294A Lecture notes