

Table of Contents

- 1 Autoencoders
 - Autoencoder Basics
- 2 Stacked Autoencoder
 - Extending Autoencoders
- 3 Sparse Autoencoder
 - k-Sparse Autoencoder
 - Sparsity Constrained Autoencoder
- 4 Denoising Autoencoder
 - Denoising Autoencoder
 - Extending Denoising Autoencoder
- 5 Conclusion
 - Advantages of Autoencoders

Autoencoder Basics


- 1 Feed-forward neural network
- 2 Used for unsupervised learning
- 3 Trained to reproduce input layer in the output layer
- 4 Training uses gradient descent

Autoencoder Basics cont.

- ① Similar to PCA*, but more flexible
- ② Autoencoders can accommodate non-linear transformations
 - 1 Consequence of flexible activation functions
 - 2 Turns out to be hugely useful

* Actually, with squared error loss and no sigmoid transformation, it is equivalent to PCA (Baldi & Hornik, 1989)

Simple Autoencoder



Simple Autoencoder cont.

From input layer to hidden layer is “encoder”:

$$f(x) = \sigma(Wx + b)$$

Hidden layer to the output is the “decoder”:

$$g(x) = \sigma(W'x + b')$$

where σ is the activation function (often sigmoid) and W' is often W^\top (known as having “tied” weights).

Simple Autoencoder cont.

- 1 Input layer \rightarrow hidden layer (fewer units) \rightarrow output layer
- 2 Hidden layers values can be thought of as lossy compression of input layer
- 3 Hidden layer with fewer units than input is often called “bottleneck” or under-complete layer

Question

In a sense, autoencoders want to learn an approximation to the identity function.

Why would this be useful?

Answer


This turns out to have several practical uses.

- 1 Autoencoders learn a kind of latent representation of the input
- 2 Can be used for initializing weights and biases prior to fitting neural net
- 3 Very well suited to dimensionality reduction in NLP tasks

Stacked Autoencoder

- 1 Sometimes called deep autoencoder
- 2 Feed-forward neural network
- 3 Has additional layers
- 4 Still “unsupervised” in the sense of no labels or real-valued outcome variable
- 5 Output layer is still \hat{x}
- 6 Trained in stages

Stacked Autoencoder



Stacked Autoencoder cont.

- ① Beneficial for data with hierarchical organization
- ② Can be difficult to train using back propagation

Sparse Autoencoder

- 1 Different approach to “compressing” representation
- 2 Hidden layer has more units than input (i.e., over-complete)
- 3 Compression is therefore achieved through sparsity


Sparse Autoencoder cont.

- 1 Several methods of introducing sparsity, we'll discuss 2
 - 1 k -sparse autoencoders (Makhzani et al., 2013)
 - 2 Constrained activation

k -sparse Autoencoder

- 1 Select k highest-activation units
- 2 Set others to 0
- 3 Compression is achieved through sparsity of model, not reduced dimensionality

k -sparse Autoencoder



Sparsity Constraint

- 1 With sigmoidal activation function, $\sigma_j(x) \in [0, 1]$
- 2 We want to constrain them to be mostly 0
- 3 Take mean activation

$$\hat{\rho}_j = \frac{1}{n} \sum_{i=1}^n [\sigma_j(x_i)]$$

such that

$$\hat{\rho}_j = \rho$$

where ρ is a hyperparameter set near 0, commonly 0.05.

Sparsity Constraint cont.

- 1 Add penalty term to loss function to achieve this sparsity constraint

$$\sum_{j=1}^m \rho \log \frac{\rho}{\hat{\rho}_j} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j}$$

- 2 This is based on Kullback-Leibler Divergence, and often written as

$$\sum_{j=1}^m \text{KL}(\rho || \hat{\rho}_j)$$

Sparsity Constraint cont.

- 1 Penalty term has the property that $\text{KL}(\rho || \hat{\rho}_j) = 0$ when $\hat{\rho}_j = \rho$
- 2 Otherwise, monotone increasing as $\hat{\rho}_j$ diverges from ρ
- 3 So our loss function is now

$$L_{\text{sparse}}(x, \hat{x}) = L(x, \hat{x}) + \lambda \sum_{j=1}^m \text{KL}(\rho || \hat{\rho}_j)$$

where $L(x, \hat{x})$ is our unconstrained loss function and λ controls the weight of sparsity penalty.

Denoising Autoencoder cont.


- ① Stochastic version of autoencoder
- ② Random noise injected into the input layer
 - 1 *Masking noise*: some fraction of elements* in x set to 0
 - 2 *Gaussian additive noise*: $\tilde{x}|x \sim \mathcal{N}(x, \sigma^2 I)$
 - 3 *Salt-and-pepper noise*: fraction of elements* in x set to either min or max possible value (often 0 or 1) according to coin toss

* chosen at random for each example

Denoising Autoencoder

- ① Output layer is evaluated against original (uncorrupted) data
- ② Goal is slightly different from simple autoencoder
 - 1 Not attempting to learn identity function
 - 2 We want robustness of representation
 - 3 Learned representation is insensitive to perturbations in the input

Denoising Autoencoder



Denoising Autoencoder cont.

- ① Training proceeds fairly similarly to standard autoencoder
- ② Crucial difference is the loss function is calculated using *uncorrupted* input layer

Extending Denoising Autoencoder

- 1 Some noise types only corrupt subset of the input's components
 - 1 Masking noise
 - 2 Salt-and-pepper noise
- 2 For these, we can extend denoising by adding emphasis on corrupted components.
- 3 Use hyperparameters α and β to control emphasis; for squared error loss, this yields

$$L(x, \hat{x}) = \alpha \left(\sum_{i \in \mathcal{I}(\tilde{x})} (x_i - \hat{x}_i)^2 \right) + \beta \left(\sum_{i \notin \mathcal{I}(\tilde{x})} (x_i - \hat{x}_i)^2 \right)$$

where $\mathcal{I}(\tilde{x})$ denotes indexes components of x that were corrupted.

Some Advantages of Autoencoders

- ① Compressed representation of features used for initializing weights
 - 1 Faster learning
 - 2 Less likely to get stuck in local minima
 - 3 Better approximation → increased accuracy
- ② Well suited to dimensionality reduction in NLP problems
- ③ Allows us to make use of unlabeled data

Disadvantages of Autoencoders

- 1 Computational burden
- 2 Limitations of unsupervised learning

Not Discussed

- 1 Winner-Take-All autoencoders (for sparsity)
- 2 Convolutional autoencoders (several varieties)
- 3 Variational autoencoders

References

- 1 Vincent P., Larochelle, H., Bengio Y., Manzagol P.A. (2008) *Extracting and Composing Robust Features with Denoising Autoencoders*.
- 2 Vincent P., Larochelle, H., LaJoie I., Bengio Y., Manzagol P.A. (2010) *Stacked Denoising Autoencoders: Learning Useful Representations in a Deep Network with a Local Denoising Criterion*
- 3 Makhzani A., Frey B. (2015) *Winner-Take-All Autoencoders*
- 4 Ng A. *Sparse Autoencoders* CS294A Lecture notes